

Gaussian processes & non-Gaussian likelihoods



ST John


ti.john@aalto.fi
infinitcuriosity.org

Finnish Center for Artificial Intelligence
& Aalto University


Gaussian Process Summer School 2024, 10 September 2024

Not Gaussian noise







How can we model this?



SVM classification



Gaussian process regression




Gaussian process **classification**

Why don't we use regression models for classification?

- Binary classification: data set $\{x_n, y_n\}_{n=1}^N$ with $y_n \in \{0, 1\}$
- We want to model $p(y_n = +1 | x_n)$
- Why not simply use a GP regression model with labels: $y_n \in \{0, 1\}$:

$$p(y_n = +1 | x_n) = f(x_n)$$



Outline:


1. **Gaussian processes with Gaussian likelihood**
 2. What is the likelihood? Connecting observations and Gaussian process prior
 3. Non-Gaussian likelihoods: what happens to the posterior?
 4. How to approximate the intractable
 5. Comparison
-
- | | |
|----------------------------------|----------------------|
| + <i>Intuitive</i> understanding | – In-depth expertise |
| + Learning the language | – Lots of maths |

Setting the scene

Gaussian process $f(\cdot)$

Distribution over functions

Marginals are Gaussian (mean and covariance)




infinitecuriosity.org/vizgp

Gaussian process conditioned on observation

Distribution over functions

Marginals are Gaussian (mean and covariance)




infinitecuriosity.org/vizgp

exact conditioning

Without noise model, we interpolate observations:


$$y(x) = f(x) + \epsilon, \quad \epsilon \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_{\text{noise}}^2)$$
$$p(y|f) = \mathcal{N}(y|f, \sigma_{\text{noise}}^2)$$



exact conditioning

Without noise model, we interpolate observations:


$$y(x) = f(x) + \epsilon, \quad \epsilon \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_{\text{noise}}^2)$$
$$p(y|f) = \mathcal{N}(y|f, \sigma_{\text{noise}}^2)$$



Gaussian noise model

Gaussian additive noise model, written two ways:


$$y(x) = f(x) + \epsilon, \quad \epsilon \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_{\text{noise}}^2)$$
$$p(y|f) = \mathcal{N}(y|f, \sigma_{\text{noise}}^2)$$



misspecified Gaussian noise model


Gaussian additive noise model, written two ways:

$$y(x) = f(x) + \epsilon, \quad \epsilon \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_{\text{noise}}^2)$$
$$p(y|f) = \mathcal{N}(y|f, \sigma_{\text{noise}}^2)$$



heavy-tailed noise model


$$y(x) = f(x) + \epsilon, \quad \epsilon \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_{\text{noise}}^2)$$
$$p(y|f) = \mathcal{N}(y|f, \sigma_{\text{noise}}^2)$$



- ✓ Gaussian processes with Gaussian likelihood
- 2. **What is the likelihood? Connecting observations and Gaussian process prior**
- 3. Non-Gaussian likelihoods: what happens to the posterior?
- 4. How to approximate the intractable
- 5. Comparison

Likelihood

Non-Gaussian observations



latent functional relationship (correlations!)
 $p(y_n | f(x_n))$

Likelihood

$$p(\mathbf{y} | \mathbf{f}) = \prod_{n=1}^N p(y_n | f_n); \quad f_n = f(x_n)$$


factorizing

Let's consider the individual (marginal, 1D) likelihood term:


$$p(y | f)$$

Function of two arguments:




$$y \mapsto p(y | f), \quad f \mapsto p(y | f)$$






$p(y|f)$: Gaussian






$p(y|f)$: Gaussian







$p(y|f)$: Bernoulli



$p(y|f)$: Beta







Two important aspects of likelihoods:

1. link functions
2. log-concavity


Link functions

$$\mathbb{E}[y] = \theta \in (0 \dots \infty)$$

$$f \sim \mathcal{N} \quad \in (-\infty \dots \infty)$$

$$\text{link}(\theta) = f$$

$$\theta = \text{invlink}(f)$$




Link functions

$$\mathbb{E}[y] = \theta \in (0 \dots \infty)$$

$$f \sim \mathcal{N} \quad \in (-\infty \dots \infty)$$

$$\text{link}(\theta) = f$$

$$\theta = \text{invlink}(f)$$




Link functions

$$\mathbb{E}[y] = \theta \in (0 \dots \infty)$$


$$f \sim \mathcal{N} \quad \in (-\infty \dots \infty)$$

$$\text{link}(\theta) = f$$

$$\theta = \text{invlink}(f)$$




(Log-)concavity




$$f(\alpha x + (1 - \alpha)y) \geq \alpha f(x) + (1 - \alpha)f(y)$$


Log-concavity of likelihoods




Functional prior $p(f)$




Functional prior $p(f)$




Functional prior $p(f)$




Joint (generative) model: $p(y, f) = p(y | f)p(f)$




Joint (generative) model: $p(y, f) = p(y | f)p(f)$




Joint (generative) model: $p(y, f) = p(y | f)p(f)$




Posterior: $p(f | y) = p(y | f)p(f) / p(y)$
 $f(x) | y(x) ?$



Posterior: $p(f | y) = p(y | f)p(f) / p(y)$




Posterior: $p(f | y) = p(y | f)p(f) / p(y)$



Back to GPs...

Posterior: $p(f | y) = p(y | f)p(f) / p(y)$ for more data



- ✓ Gaussian processes with Gaussian likelihood
- ✓ What is the likelihood? Connecting observations and Gaussian process prior
- 3. **Non-Gaussian likelihoods: what happens to the posterior?**
- 4. How to approximate the intractable
- 5. Comparison

Posterior

Likelihood

$$p(\mathbf{y} \mid \mathbf{f})$$

Joint distribution

$$p(\mathbf{y}, \mathbf{f}) = p(\mathbf{y} \mid \mathbf{f})p(\mathbf{f})$$

Posterior

$$\mathbf{f} \mapsto p(\mathbf{f} \mid \mathbf{y}) = \frac{p(\mathbf{y} \mid \mathbf{f})p(\mathbf{f})}{p(\mathbf{y})}$$

$$\mathbf{y} \mapsto (\mathbf{f} \mapsto p(\mathbf{f} \mid \mathbf{y}))$$

Posterior predictions

At new point x^* :

$$p(f^* | x^*, \mathbf{x}, \mathbf{y}) = \int p(f^* | x^*, \mathbf{x}, \mathbf{f}) p(\mathbf{f} | \mathbf{x}, \mathbf{y}) d\mathbf{f}$$

At training data:

$$p(\mathbf{f} | \mathbf{x}, \mathbf{y}) = \frac{p(\mathbf{f} | \mathbf{x}) \prod_{n=1}^N p(y_n | f(x_n))}{\int p(\mathbf{f}' | \mathbf{x}) \prod_{n=1}^N p(y_n | f'(x_n)) d\mathbf{f}'}$$

$$p(\mathbf{f} | \mathbf{y}) = \frac{1}{Z} p(\mathbf{f}) \prod_{n=1}^N p(y_n | f_n)$$

$$Z = p(\mathbf{y} | \mathcal{M}) = \int p(\mathbf{f} | \mathcal{M}) \prod_{n=1}^N p(y_n | f_n, \mathcal{M}) d\mathbf{f}$$

“marginal likelihood” or “evidence” given **model** $\mathcal{M}$

$$p(\mathbf{f} | \mathbf{y}) = \frac{1}{Z} p(\mathbf{f}) \prod_{n=1}^N p(y_n | f_n)$$


Gaussian (process) prior $p(f(\cdot)) \dots$ $p(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mathbf{0}, \mathbf{K})$

& Gaussian likelihood: conjugate case $\rightarrow$ posterior Gaussian


& **non**-Gaussian $p(y|f)$ $\rightarrow p(\mathbf{f} | \mathbf{y})$ also **non**-Gaussian, **intractable**

1D examples


Gaussian



Student's t




Bernoulli




Bernoulli example in 2D


Bernoulli





Prior



Likelihood



Posterior



Posterior for N observations

$$p(\mathbf{f} | \mathbf{y}) = \frac{p(\mathbf{f}) \prod_{n=1}^N p(y_n | f_n)}{\int p(\mathbf{f}') \prod_{n=1}^N p(y_n | f'_n) d\mathbf{f}'}$$

$$f_1 = f(x_1)$$

$$f_2 = f(x_2)$$

$$\vdots$$

$$f_N = f(x_N)$$

Summary so far

- What is the likelihood $p(y|f)$?
- When is it non-Gaussian?
- Why does the posterior $p(f|y)$ become intractable?

Questions?! :)

- ✓ Gaussian processes with Gaussian likelihood
- ✓ What is the likelihood? Connecting observations and Gaussian process prior
- ✓ Non-Gaussian likelihoods: what happens to the posterior?
- 4. **How to approximate the intractable**
- 5. Comparison

Approximations

■ Joint model:

$$p(\mathbf{y}, \mathbf{f}) = p(\mathbf{y} | \mathbf{f}) p(\mathbf{f}) = \prod_{n=1}^N p(y_n | f_n) \mathcal{N}(\mathbf{f} | \mathbf{0}, \mathbf{K})$$

■ Posterior distribution at training points:

$$p(\mathbf{f} | \mathbf{y}) = \frac{p(\mathbf{y} | \mathbf{f}) p(\mathbf{f})}{p(\mathbf{y})} \approx q(\mathbf{f})$$

■ Posterior of f^* for new test point $\mathbf{x}^*$:

$$p(f^* | \mathbf{y}) = \int p(f^* | \mathbf{f}) p(\mathbf{f} | \mathbf{y}) d\mathbf{f} \approx \int p(f^* | \mathbf{f}) q(\mathbf{f}) d\mathbf{f} \equiv q(f^*)$$


■ Predictive distribution

$$p(y^* | \mathbf{y}) = \int p(y^* | f^*) p(f^* | \mathbf{y}) df^* \approx \int p(y^* | f^*) q(f^*) df^*$$

Analytically intractable distributions!


Approximating distributions

- Delta distribution
 - ▶ Point estimate
- Mixture of delta distributions
 - ▶ **Markov Chain Monte Carlo (MCMC)**
 - ▶ Neural network ensembles...
- Gaussian distribution
 - ▶ Laplace Approximation (LA)
 - ▶ Variational Bayes/Variational Inference (VB / VI)
 - ▶ Expectation Propagation (EP), PowerEP, ...
- Mixture of Gaussians
- ...



Markov Chain Monte Carlo

Markov Chain




- Samples $x_1, \dots, x_T$
- “Markov” = 1-step history
- $x_{t+1} \sim p(x_{t+1} | x_t)$, independent of $x_{t-1}, \dots, x_1$

Markov Chain Monte Carlo (MCMC)

Generate samples $\{x_t\} \sim p(f | y)$

Requires:

- *unnormalized* posterior $h(f) = p(y|f)p(f)$
- Markov proposal $q(x' | x_t)$
- initial x_0



In each iteration t :

1. Random proposal $x' \sim q(x' | x_t)$
2. Acceptance probability $\frac{h(x')}{h(x_t)} \rightarrow$ ensures sampling from $p(f | y)$


accept: $x_{t+1} = x'$ reject: copy $x_{t+1} = x_t$

$h(x') > h(x_t)$: always accepts $\rightarrow$ climbs uphill

Demo: MCMC in 2D

tinyurl.com/nongaussian-inference-viz-v1

Marginals



MCMC: important properties

- burn-in
- acceptance ratio
- auto-correlation, effective sample size (ESS); thinning to save memory
- mixing and multiple chains ($\hat{R}$)
- better proposals (HMC, NUTS) → use robust implementations
 - + very accurate (gold-standard)
 - very slow, predictions require keeping all (thinned) samples around

Michael Betancourt's betanalpha.github.io/writing/

MCMC: robust implementations

■ Stan



■ PyMC3



PYMC3

■ Pyro & NumPyro




■ TensorFlow Probability (GPflow)



GPflow

■ Turing.jl



- ✓ Gaussian processes with Gaussian likelihood
- ✓ What is the likelihood? Connecting observations and Gaussian process prior
- ✓ Non-Gaussian likelihoods: what happens to the posterior?

4. **How to approximate the intractable**

- ✓ with samples: MCMC

4.2 **with Gaussians**

- Laplace Approximation
- Variational Inference
- Expectation Propagation

5. Comparison

Gaussian approximations

Approximating the exact posterior with Gaussian

Approximating the posterior at observations:

$$p(\mathbf{f} | \mathbf{y}) \approx q(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mu = ?, \Sigma = ?)$$

Predictions at new points:

$$p(f^* | x^*, \mathbf{y}) \approx q(f^*) = \int p(f^* | x^*, \mathbf{f}) q(\mathbf{f}) \mathrm{d}\mathbf{f}$$

Demo: What does this mean for Gaussian processes?

tinyurl.com/nongaussian-inference-viz-v1

Choosing μ and Σ for $q(\mathbf{f})$

$$p(\mathbf{f} | \mathbf{y}) \approx q(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mu = ?, \Sigma = ?)$$

locally: match mean &
variance at point

globally: minimise divergence

**Laplace
approximation**

Variational
Inference (VI)

Expectation
Propagation (EP)

Laplace approximation

Laplace approximation

Idea: log of Gaussian pdf = quadratic polynomial

$$p_{\mathcal{N}}(\mathbf{f}) = \frac{1}{\sqrt{(2\pi)^k |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{f} - \mu)^{\top} \Sigma^{-1} (\mathbf{f} - \mu)\right)$$


quadratic polynomial through approximation:

2nd-order Taylor expansion of log of $h(f) = p(y|f)p(f)$ at $\hat{f}$


$$g(x + \delta) \approx g(x) + \left(\frac{dg}{dx}(x)\right)\delta + \frac{1}{2!}\left(\frac{d^2g}{dx^2}(x)\right)\delta^2$$

1. Find **mode** of posterior
2nd-order gradient optimisation (e.g. Newton's method)
2. Match **curvature** (Hessian) at mode


$$p(f | y) = \frac{1}{Z} p(y | f) p(f)$$




$$\log p(f | y) = -\log Z + \log p(y | f) + \log p(f)$$




$$\log p(f | y) = -\log Z + \log h(f)$$




Newton's method




Newton's method




Newton's method




Newton's method




Newton's method




$$\log p(f | y) + \log Z = \log h(f) \approx \mathcal{O}(f^2)$$




$$p(f | y) Z \approx \exp(\mathcal{O}(f^2))$$




$$p(f | y) \approx \mathcal{N}(f | \hat{f}, -(\mathrm{d}^2 \log h / \mathrm{d}f^2)^{-1})$$




$$p(f | y) \approx \mathcal{N}(f | \hat{f}, -(\mathrm{d}^2 \log h / \mathrm{d}f^2)^{-1}) = q(f)$$



Laplace in 2D example



Laplace in 2D: marginals



Laplace approximation: important properties

- find mode: Newton's method
- match curvature (Hessian) at mode
- “point estimate++”
 - + simple, fast
 - poor approximation if mode is not representative (e.g. Bernoulli)
 - may not converge for non-log-concave likelihoods [1]

Choosing μ and Σ for $q(\mathbf{f})$

$$p(\mathbf{f} | \mathbf{y}) \approx q(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mu = ?, \Sigma = ?)$$

locally: match mean &
variance at point

globally: minimise divergence

Laplace
approximation

Variational
Inference (VI)

Expectation
Propagation (EP)

Variational Bayes (VB)

Variational Inference (VI)

Variational inference: the big picture

Recipe for approximating intractable distribution

$p \in \mathcal{P}$


1. Define some “simple” family of distributions $\mathcal{Q}$.
2. Define some way to compute a “distance” $\mathbb{D}[p, q]$ between intractable distribution p and each distribution $q \in \mathcal{Q}$

$$\mathbb{D}[p, q_1] > \mathbb{D}[p, q_2]$$

3. Search for $q \in \mathcal{Q}$ such that $\mathbb{D}[p, q]$ is minimized

$$q^* = \arg \min_{q \in \mathcal{Q}} \mathbb{D}[p, q]$$

4. Use q^* as an approximation of p



How to “measure distances” between distributions?

Here: *Kullback–Leibler (KL) divergence*

$$\mathbb{D}[p, q] := \text{KL}[q \parallel p] = \int q(f) \log \frac{q(f)}{p(f)} \mathrm{d}f = \mathbb{E}_q \left[\log \frac{q(f)}{p(f)} \right]$$

Important properties:

1. Non-symmetric: $\text{KL}[q \parallel p] \neq \text{KL}[p \parallel q]$
2. Positive: $\text{KL} \geq 0$ (Gibbs’ inequality)
3. Minimum: $\text{KL}[q \parallel p] = 0 \iff q \equiv p$.

$$q(\mathbf{f}) = \mathcal{N}(\mu, \Sigma)$$

$$\operatorname{argmin}_{\mu, \Sigma} \operatorname{KL} [q(\mathbf{f}) \| p(\mathbf{f} | \mathbf{y})]$$

Variational inference: Minimizing $\text{KL}[q(\mathbf{f}) \| p(\mathbf{f} | \mathbf{y})]$

$$\begin{aligned}\text{KL}[q(\mathbf{f}) \| p(\mathbf{f} | \mathbf{y})] &= \int q(\mathbf{f}) \left[\log \frac{q(\mathbf{f})}{p(\mathbf{f} | \mathbf{y})} \right] d\mathbf{f} \\&= \int q(\mathbf{f}) [\log q(\mathbf{f}) - \log p(\mathbf{f} | \mathbf{y})] d\mathbf{f} \\&= \int q(\mathbf{f}) \left[\underbrace{\log q(\mathbf{f}) - \log p(\mathbf{f})}_{\text{KL}[q(\mathbf{f}) \| p(\mathbf{f})]} - \log p(\mathbf{y} | \mathbf{f}) + \log p(\mathbf{y}) \right] d\mathbf{f} \\&= \int q(\mathbf{f}) \left[\log \frac{q(\mathbf{f})}{p(\mathbf{f})} \right] d\mathbf{f} - \int q(\mathbf{f}) [\log p(\mathbf{y} | \mathbf{f})] d\mathbf{f} + \log p(\mathbf{y}) \\&= \text{KL}[q(\mathbf{f}) \| p(\mathbf{f})] - \int q(\mathbf{f}) [\log p(\mathbf{y} | \mathbf{f})] d\mathbf{f} + \log p(\mathbf{y}) \\ \log p(\mathbf{y}) &= \int q(\mathbf{f}) [\log p(\mathbf{y} | \mathbf{f})] d\mathbf{f} - \text{KL}[q(\mathbf{f}) \| p(\mathbf{f})] + \text{KL}[q(\mathbf{f}) \| p(\mathbf{f} | \mathbf{y})]\end{aligned}$$

Variational inference: Minimizing $\text{KL}[q(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})]$ by bounding

$$\begin{aligned}\log p(\mathbf{y}) &= \underbrace{\int q(\mathbf{f}) [\log p(\mathbf{y} \mid \mathbf{f})] d\mathbf{f} - \text{KL}[q(\mathbf{f}) \parallel p(\mathbf{f})]}_{\mathcal{L}[q]} + \underbrace{\text{KL}[q(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})]}_{\geq 0} \\ &\geq \int q(\mathbf{f}) [\log p(\mathbf{y} \mid \mathbf{f})] d\mathbf{f} - \text{KL}[q(\mathbf{f}) \parallel p(\mathbf{f})] = \mathcal{L}[q]\end{aligned}$$

- $\log p(\mathbf{y})$ is a constant
- $\mathcal{L}[q]$ does **not** depend on $p(\mathbf{f} \mid \mathbf{y})$
- $\mathcal{L}[q] \leq \log p(\mathbf{y})$, so $\mathcal{L}[q]$ is *lower bound* on marginal log likelihood $\log p(\mathbf{y})$
- Maximizing $\mathcal{L}[q]$ is equivalent to minimizing $\text{KL}[q(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})]$

Key take-away: we can fit variational approximation q by optimizing $\mathcal{L}$

$$\log p(\mathbf{y}) \geq \mathcal{L}[q] = \underbrace{\int q(\mathbf{f}) [\log p(\mathbf{y} | \mathbf{f})] d\mathbf{f}}_{\text{data fit}} - \underbrace{\text{KL}[q(\mathbf{f}) \parallel p(\mathbf{f})]}_{\text{regularization}}$$

$\mathcal{L}[q]$ often called the *Evidence Lower Bound* (ELBO)

- We approximate $p(\mathbf{f} | \mathbf{y}) \approx q(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mu = ?, \Sigma = ?)$
- Defining $\lambda = \{\mu, \Sigma\}$, we can write $\mathcal{L}[q] = \mathcal{L}(\lambda)$
- In practice, we optimize $\mathcal{L}(\lambda)$ using gradient-based methods

Likelihood term (data fit)

Integral separates for a factorizing likelihood:

$$\begin{aligned} & \int q(\mathbf{f}) [\log p(\mathbf{y} | \mathbf{f})] d\mathbf{f} \\ &= \sum_{n=1}^N \int q(f_n) [\log p(y_n | f_n)] df_n \end{aligned}$$

Sum over 1D integrals:


- analytic for some (e.g. Exponential, Gamma, Poisson)
- fast and accurate to approximate numerically (for example Gauss–Hermite quadrature)
- Monte Carlo (e.g. multi-class classification)

Take away #2: We can tractably optimize the bound for non-Gaussian likelihoods

1D Toy example


Assume model $p(y, f)$ with some intractable posterior $p(f | y)$

- In 1D: $\mathcal{Q}$ is the set of univariate Gaussians, i.e. $q_\lambda(f) = \mathcal{N}(f | m, v)$, and $\lambda = \{m, v\}$
- Initialization: $q(f) = \mathcal{N}(f | 0, 1)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$




1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$



1D Toy example

- Gradient ascent: $\lambda_{i+1} = \lambda_i + \eta \nabla_{\lambda} \mathcal{L}(\lambda)$
- $\log p(\mathbf{y}) = \mathcal{L}(\lambda) + \mathbb{D}[q_{\lambda}(\mathbf{f}) \parallel p(\mathbf{f} \mid \mathbf{y})] \geq \mathcal{L}(\lambda)$



Variational inference: important properties

- principled: directly minimising divergence from true posterior
- mode-seeking (e.g. multi-modal posterior: fits just one, if q is unimodal)
 - + minimises a true lower bound $\rightarrow$ convergence
 - underestimates variance

Minimising divergences

$$p(\mathbf{f} | \mathbf{y}) \approx q(\mathbf{f}) = \mathcal{N}(\mathbf{f} | \mu = ?, \Sigma = ?)$$

✓ $\min \text{KL}[q(\mathbf{f}) \| p(\mathbf{f} | \mathbf{y})]$: Variational Inference

2. $\min \text{KL}[p(\mathbf{f} | \mathbf{y}) \| q(\mathbf{f})]$: **Expectation Propagation**

Expectation Propagation (EP)

Expectation Propagation

Can we minimise KL divergence in opposite direction?

$$q(\mathbf{f}) = \operatorname{argmin}_{\mu, \Sigma} \operatorname{KL} [p(\mathbf{f} | \mathbf{y}) \| q(\mathbf{f})] = \operatorname{argmin}_{\mu, \Sigma} \int p(\mathbf{f} | \mathbf{y}) \left[\log \frac{p(\mathbf{f} | \mathbf{y})}{q(\mathbf{f})} \right] d\mathbf{f}$$

Exact posterior: $p(\mathbf{f} | \mathbf{y}) \propto p(\mathbf{f}) \prod_{n=1}^N p(y_n | f_n)$

Approximate posterior: $q(\mathbf{f}) \propto p(\mathbf{f}) \prod_{n=1}^N t_n(f_n)$

$$t_n = Z_n \mathcal{N}(f_n | \tilde{\mu}_n, \tilde{\sigma}_n^2) \quad \text{“sites”}$$


- Expectation propagation iteratively updates the sites for each data point
- minimizes KL of local approximations to the posterior:


$$\min \operatorname{KL} [q(f_n) \frac{p(y_n | f_n)}{t_n(f_n)} \| q(f_n) \frac{t'_n(f_n)}{t_n(f_n)}]$$

Expectation Propagation: important properties

- multiple passes required to converge
- moment-matching (e.g. covering multiple modes)
 - + effective for classification
 - not guaranteed to converge
 - updates may be invalid (non-log-concave likelihoods) [2]

Comparison 2D





- ✓ Gaussian processes with Gaussian likelihood
- ✓ What is the likelihood? Connecting observations and Gaussian process prior
- ✓ Non-Gaussian likelihoods: what happens to the posterior?
- ✓ How to approximate the intractable
 - ✓ with samples: MCMC
 - ✓ with Gaussians
 - Laplace
 - Variational Inference
 - Expectation Propagation

5. Comparison

Comparison

Comparison

MCMC

- ▶ samples
- ▶ gold standard
- ▶ slow

Laplace

- ▶ $\mathcal{N}$ = curvature at mode
- ▶ simple & fast
- ▶ often poor approximation

Variational Inference

- ▶ $\mathcal{N}$ minimises $\text{KL}[q(\mathbf{f})||p(\mathbf{f} | \mathbf{y})]$
- ▶ principled, any likelihood
- ▶ underestimates variance

Expectation Propagation

- ▶ $\mathcal{N}$ matches marginal moments
- ▶ good calibration in classification
- ▶ may not converge


What we did not cover...

- More complex likelihoods (heteroskedastic, zero-inflated, multi-stage...)
- Marginal likelihood approximations for hyperparameter learning [3]
- How parametrisation affects Gaussianity of $p(\mathbf{f} | \mathbf{y})$
- Connections between EP and VI (“PowerEP”, CVI dual parameterization) [4, 5]
- Other divergences, generalised VI, ...
- Combinations of MCMC and variational methods
- Augmenting likelihood with auxiliary variable
→ conditionally conjugate model [6]

Take-away

We can...

- create **richer models** with likelihoods beyond the Gaussian
- **learn latent functions** that form the connection between data points
- handle the non-Gaussian posterior with **approximations**
- **trade off** speed, accuracy, and ease-of-use



References I



Marcelo Hartmann and Jarno Vanhatalo.

Laplace approximation and natural gradient for Gaussian process regression with heteroscedastic student- t model.

Statistics and Computing, 29(4):753–773, October 2018.



Pasi Jylänki, Jarno Vanhatalo, and Aki Vehtari.

Robust Gaussian process regression with a student- t likelihood.

Journal of Machine Learning Research, 12(99):3227–3257, 2011.



Hannes Nickisch and Carl Edward Rasmussen.

Approximations for binary Gaussian process classification.

Journal of Machine Learning Research, 9(67):2035–2078, 2008.



Thang D. Bui, Josiah Yan, and Richard E. Turner.

A unifying framework for Gaussian process pseudo-point approximations using Power Expectation Propagation.

Journal of Machine Learning Research, 18(104):1–72, 2017.

References II



Vincent Adam, Paul Chang, Mohammad Emtiyaz E Khan, and Arno Solin.

Dual parameterization of sparse variational gaussian processes.

In M. Ranzato, A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, volume 34, pages 11474–11486, 2021.



Théo Galy-Fajou, Florian Wenzel, and Manfred Opper.

Automated augmented conjugate inference for non-conjugate Gaussian process models, 2020.



James Hensman, Nicolo Fusi, and Neil D. Lawrence.

Gaussian processes for big data.

UAI, 2013.



Malte Kuss and Carl Edward Rasmussen.

Assessing approximate inference for binary Gaussian process classification.

Journal of Machine Learning Research, 6(57):1679–1704, 2005.



Alan Saul.

Gaussian process based approaches for survival analysis, 2017.

References III



Aki Vehtari, Andrew Gelman, Tuomas Sivula, Pasi Jylänki, Dustin Tran, Swupnil Sahai, Paul Blomstedt, John P. Cunningham, David Schiminovich, and Christian P. Robert.
Expectation Propagation as a way of life: A framework for Bayesian inference on partitioned data.
Journal of Machine Learning Research, 21(17):1–53, 2020.



Will Penny.
Bayesian inference course: Variational inference, 2013.